

Applied Econometrics with R

Karim Kilani

2019

Settings and appearance

```
knitr::opts_chunk$set(echo=TRUE)
options(prompt="R>", digits=4)
1.32222222
```

```
## [1] 1.322
```

- Replaces R prompt `>` by a new prompt `R>`.
- Reduces the number of digits shown when printing numbers from 7 (default) to 4 digits.

Introductory R Session

Example 1: The demand for economics journals

```
data("Journals", package="AER")
dim(Journals)
## [1] 180 10


names(Journals)[1:4]
## [1] "title"      "publisher"  "society"    "price"

names(Journals)[5:7]
## [1] "pages"      "charpp"     "citations"

names(Journals)[8:10]
## [1] "foundingyear" "subs"       "field"
```


- We study the relation between the demand for economics journals and their price.
- Demand is measured by the number of library subscriptions (subs).
- Price is measured by is the price per citation.
- A scatterplot of the logarithms of the variables can be constructed.

```
plot(log(subs)~log(price/citations),data=Journals)
```



- The number of subscriptions is decreasing with price.
- The corresponding linear regression model can be easily fitted by ordinary least squares (OLS).

```
j_lm<-lm(log(subs)~log(price/citations),data=Journals)
plot(log(subs)~log(price/citations),data=Journals)
abline(j_lm,col="red")
```



- A detailed summary of the fitted model can be obtained.


```
summary(j_lm)
```

```
##
```

```
## Call:
```

```
## lm(formula = log(subs) ~ log(price/citations), data = Journa
```

```
##
```

```
## Residuals:
```

```
##      Min       1Q   Median       3Q      Max
```

```
## -2.7248 -0.5361  0.0372  0.4662  1.8481
```

```
##
```

```
## Coefficients:
```

```
##              Estimate Std. Error t value Pr(>|t|)
```

```
## (Intercept)          4.7662     0.0559   85.2  <2e-16 *
```

```
## log(price/citations) -0.5331     0.0356  -15.0  <2e-16 *
```

```
## ---
```

```
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
##
```

```
## Residual standard error: 0.75 on 178 degrees of freedom
```

```
## Multiple R-squared:  0.557    Adjusted R-squared:  0.555
```

- The estimated elasticity of the demand with respect to the price per citation is -0.5331 , which is significantly different from 0.
- The $R^2 = 0.557$ of the model is quite satisfactory.

Example 2: Determinants of wages

- The application is the estimation of a wage equation in semi-logarithmic form based on data taken from Berndt (1991).
- They represent a random subsample (533 observations) of cross-section data from the May 1985 Current Population Survey.

```
data("CPS1985",package="AER")
cps<-CPS1985
head(cps)
```

```
##      wage education experience age ethnicity region gender
## 1      5.10         8         21  35  hispanic  other female
## 1100  4.95         9         42  57      cauc  other female
## 2      6.67        12          1  19      cauc  other  male
## 3      4.00        12          4  22      cauc  other  male
## 4      7.50        12         17  35      cauc  other  male
## 5     13.07        13          9  28      cauc  other  male
##
##      sector union married
## 1  manufacturing    no    yes
## 1100 manufacturing    no    yes
## 2  manufacturing    no    no
## 3      other        no    no
## 4      other        no    yes
## 5      other        yes    no
```

- A wage equation is estimated with $\log(\text{wage})$ as the dependent variable and education and experience (in number of years) as regressors.
- We estimate a multiple linear regression model by OLS (again via `lm()`).
- Quantile regressions (a refinement of the regression model) are fitted using the function `rq()` from the package `quantreg`.
- We need to specify τ , the quantiles that are to be modeled.
- We set (0.2, 0.35, 0.5, 0.65, 0.8).

```
library("quantreg")
```

```
## Loading required package: SparseM
```

```
##
```

```
## Attaching package: 'SparseM'
```

```
## The following object is masked from 'package:base':
```

```
##
```

```
##      backsolve
```

```
cps_lm <- lm(log(wage) ~ experience + I(experience^2) +  
+ education, data = cps)  
cps_rq <- rq(log(wage) ~ experience + I(experience^2) +  
+ education, data = cps, tau = seq(0.2, 0.8, by = 0.15))
```

```
cps2 <- data.frame(education = mean(cps$education),  
  experience = min(cps$experience):max(cps$experience))  
head(cps2)
```

```
##      education experience  
## 1         13.02          0  
## 2         13.02          1  
## 3         13.02          2  
## 4         13.02          3  
## 5         13.02          4  
## 6         13.02          5
```



```
cps2 <- cbind(cps2, predict(cps_lm, newdata = cps2,  
interval = "prediction"))  
head(cps2)
```

##	education	experience	fit	lwr	upr
## 1	13.02	0	1.689	0.7756	2.602
## 2	13.02	1	1.723	0.8110	2.635
## 3	13.02	2	1.757	0.8451	2.668
## 4	13.02	3	1.789	0.8781	2.700
## 5	13.02	4	1.820	0.9098	2.730
## 6	13.02	5	1.850	0.9404	2.760


```
cps2 <- cbind(cps2,predict(cps_rq, newdata = cps2, type = ""))
cps2[1:4,1:5]
```

##	education	experience	fit	lwr	upr
## 1	13.02	0	1.689	0.7756	2.602
## 2	13.02	1	1.723	0.8110	2.635
## 3	13.02	2	1.757	0.8451	2.668
## 4	13.02	3	1.789	0.8781	2.700

```
cps2[1:4,6:10]
```

##	tau= 0.20	tau= 0.35	tau= 0.50	tau= 0.65	tau= 0.80
## 1	1.310	1.424	1.597	1.766	1.988
## 2	1.347	1.471	1.641	1.809	2.026
## 3	1.382	1.516	1.684	1.850	2.064
## 4	1.415	1.559	1.725	1.890	2.101

```
plot(log(wage) ~ experience, data = cps)
for(i in 6:10) lines(cps2[,i]~experience,
data = cps2,col="red")
```



Getting Started

- <http://www.R-project.org/>

Packages

```
library("AER")
```

```
## Loading required package: car
## Loading required package: carData
## Loading required package: lmtest
## Loading required package: zoo
##
## Attaching package: 'zoo'
## The following objects are masked from 'package:base':
##
##      as.Date, as.Date.numeric
## Loading required package: sandwich
## Loading required package: survival
##
## Attaching package: 'survival'
## The following object is masked from 'package:quantreg':
```

Working with R

Handling objects

```
objects()
```

```
## [1] "cps"      "cps_lm"   "cps_rq"   "CPS1985"  "cps2"
## [7] "j_lm"     "Journals"
```

- There are currently eight objects, resulting from the introductory session
- This is not the complete list of available objects, given that some objects must already exist prior to the execution of any commands.

```
search()
```

```
## [1] ".GlobalEnv"          "package:AER"          "package:survi
## [4] "package:sandwich"    "package:lmtest"       "package:zoo"
## [7] "package:car"         "package:carData"      "package:quant
## [10] "package:SparseM"     "package:stats"        "package:graph
## [13] "package:grDevices"   "package:utils"        "package:datas
## [16] "package:methods"     "Autoloads"            "package:base"
```



```
x <- 2  
objects()
```

```
## [1] "cps"      "cps_lm"   "cps_rq"   "CPS1985"  "cps2"  
## [7] "j_lm"     "Journals" "x"
```

```
remove(x)  
objects()
```

```
## [1] "cps"      "cps_lm"   "cps_rq"   "CPS1985"  "cps2"  
## [7] "j_lm"     "Journals"
```

Calling functions

```
log(16, 2)
```

```
## [1] 4
```

```
log(x = 16, 2)
```

```
## [1] 4
```

```
log(16, base = 2)
```

```
## [1] 4
```

```
log(base = 2, x = 16)
```

```
## [1] 4
```

Classes and generic functions

```
class(CPS1985)
```

```
## [1] "data.frame"
```

```
##File management
```

```
getwd()
```

```
## [1] "/cloud/project"
```

Getting Help

- `?options` or `help("options")`.
- `example("lm")`.

R as a Calculator

```
1+1
```

```
## [1] 2
```

```
2^3
```

```
## [1] 8
```

```
log(exp(sin(pi/4)^2)*exp(cos(pi/4)^2))
```

```
## [1] 1
```

Vector arithmetic

```
x <- c(1.8, 3.14, 4, 88.169, 13)
```

```
length(x)
```

```
## [1] 5
```

```
2*x+3
```

```
## [1] 6.60 9.28 11.00 179.34 29.00
```

```
5:1*x+1:5
```

```
## [1] 10.00 14.56 15.00 180.34 18.00
```


Subsetting vectors

```
x
```

```
## [1] 1.80 3.14 4.00 88.17 13.00
```

```
x[c(1,4)]
```

```
## [1] 1.80 88.17
```

```
x[-c(2,3,5)]
```

```
## [1] 1.80 88.17
```

Patterned vectors

```
ones <- rep(1,10)
even <- seq(from = 2, to = 20, by = 2)
trend <- 1981:2005
c(ones,even)
```

```
## [1] 1 1 1 1 1 1 1 1 1 1 2 4 6 8 10 12 14 16
```

Matrix Operations

```
A <- matrix(1:6, nrow = 2)
```

```
A
```

```
##      [,1] [,2] [,3]
```

```
## [1,]    1    3    5
```

```
## [2,]    2    4    6
```

```
matrix(1:6, ncol = 3)
```

```
##      [,1] [,2] [,3]
```

```
## [1,]    1    3    5
```

```
## [2,]    2    4    6
```

Transpose

A

```
##      [,1] [,2] [,3]
## [1,]    1    3    5
## [2,]    2    4    6
```

$t(A)$

```
##      [,1] [,2]
## [1,]    1    2
## [2,]    3    4
## [3,]    5    6
```

Dimensions

```
dim(A)
```

```
## [1] 2 3
```

```
nrow(A)
```

```
## [1] 2
```

```
ncol(A)
```

```
## [1] 3
```

Inverting of a Matrix

```
A1 <- A[1:2, c(1, 3)]  
solve(A1)
```

```
##      [,1] [,2]  
## [1,] -1.5  1.25  
## [2,]  0.5 -0.25
```

```
A1 %*% solve(A1)
```

```
##      [,1] [,2]  
## [1,]    1    0  
## [2,]    0    1
```

Patterned matrices

```
diag(4)
```

```
##      [,1] [,2] [,3] [,4]
## [1,]    1    0    0    0
## [2,]    0    1    0    0
## [3,]    0    0    1    0
## [4,]    0    0    0    1
```

```
diag(1,4,4)
```

```
##      [,1] [,2] [,3] [,4]
## [1,]    1    0    0    0
## [2,]    0    1    0    0
## [3,]    0    0    1    0
## [4,]    0    0    0    1
```

```
diag(rep(c(1,2),c(10, 10)))
```

```
##      [,1] [,2] [,3] [,4] [,5] [,6] [,7] [,8] [,9] [,10] [,11] [,12]
## [1,]    1    0    0    0    0    0    0    0    0    0    0    0
## [2,]    0    1    0    0    0    0    0    0    0    0    0    0
## [3,]    0    0    1    0    0    0    0    0    0    0    0    0
## [4,]    0    0    0    1    0    0    0    0    0    0    0    0
## [5,]    0    0    0    0    1    0    0    0    0    0    0    0
## [6,]    0    0    0    0    0    1    0    0    0    0    0    0
## [7,]    0    0    0    0    0    0    1    0    0    0    0    0
## [8,]    0    0    0    0    0    0    0    1    0    0    0    0
## [9,]    0    0    0    0    0    0    0    0    1    0    0    0
## [10,]   0    0    0    0    0    0    0    0    0    1    0    0
## [11,]   0    0    0    0    0    0    0    0    0    0    1    0
## [12,]   0    0    0    0    0    0    0    0    0    0    0    1
```



```
diag(rep(c(1,2),c(5,5)))
```

##		[,1]	[,2]	[,3]	[,4]	[,5]	[,6]	[,7]	[,8]	[,9]	[,10]
##	[1,]	1	0	0	0	0	0	0	0	0	0
##	[2,]	0	1	0	0	0	0	0	0	0	0
##	[3,]	0	0	1	0	0	0	0	0	0	0
##	[4,]	0	0	0	1	0	0	0	0	0	0
##	[5,]	0	0	0	0	1	0	0	0	0	0
##	[6,]	0	0	0	0	0	2	0	0	0	0
##	[7,]	0	0	0	0	0	0	2	0	0	0
##	[8,]	0	0	0	0	0	0	0	2	0	0
##	[9,]	0	0	0	0	0	0	0	0	2	0
##	[10,]	0	0	0	0	0	0	0	0	0	2

- Yields a diagonal matrix of size 10×10 whose first 5 diagonal elements are 1, while the remaining ones are 2.

```
A1
```

```
##      [,1] [,2]  
## [1,]    1    5  
## [2,]    2    6
```

```
diag(A1)
```

```
## [1] 1 6
```

- Extract the diagonal from an existing matrix.

```
A1
```

```
##      [,1] [,2]  
## [1,]    1    5  
## [2,]    2    6
```

```
upper.tri(A1)
```

```
##      [,1] [,2]  
## [1,] FALSE TRUE  
## [2,] FALSE FALSE
```

```
lower.tri(A1)
```

```
##      [,1] [,2]  
## [1,] FALSE FALSE  
## [2,]  TRUE FALSE
```

*

Posi-
tions

- Add a column of ones to $A1$.

Combining matrices

```
rbind(A1,diag(4, 2))
```

```
##      [,1] [,2]  
## [1,]    1    5  
## [2,]    2    6  
## [3,]    4    0  
## [4,]    0    4
```

- Combines $A1$ and $\text{diag}(4, 2)$ by rows.

R as a Programming Language

The mode of a vector

```
x <- c(1.8, 3.14, 4, 88.169, 13)
```

Logical and character vectors

```
x > 3.5
```

```
## [1] FALSE FALSE TRUE TRUE TRUE
```

```
names(x) <- c("a", "b", "c", "d", "e")
```

```
x
```

```
##      a      b      c      d      e
```

```
##  1.80  3.14  4.00 88.17 13.00
```

```
names(x) <- letters[1:5]
```

- Assign labels to objects.

```
names(x) <- letters[1:5]
```

```
x
```

```
##      a      b      c      d      e  
## 1.80  3.14  4.00 88.17 13.00
```

More on subsetting

```
x[3:5]
```

```
##      c      d      e  
## 4.00 88.17 13.00
```

```
x[c("c", "d", "e")]
```

```
##      c      d      e  
## 4.00 88.17 13.00
```

```
x[x > 3.5]
```

```
##      c      d      e  
## 4.00 88.17 13.00
```

Lists

```
mylist <- list(sample=rnorm(5),family="normal distribution",  
parameters=list(mean=0,sd=1))
```

```
mylist
```

```
## $sample
## [1]  1.1241 -0.3772 -0.4719 -0.3263 -0.9550
##
## $family
## [1] "normal distribution"
##
## $parameters
## $parameters$mean
## [1] 0
##
## $parameters$sd
## [1] 1
```

```
mylist[[1]]
```

```
## [1] 1.1241 -0.3772 -0.4719 -0.3263 -0.9550
```

```
mylist[["sample"]]
```

```
## [1] 1.1241 -0.3772 -0.4719 -0.3263 -0.9550
```

```
mylist$sample
```

```
## [1] 1.1241 -0.3772 -0.4719 -0.3263 -0.9550
```



```
mylist[[3]]$mean
```

```
## [1] 0
```

```
mylist[[3]]$sd
```

```
## [1] 1
```

Logical comparisons

```
x <- c(1.8, 3.14, 4, 88.169, 13)
```

```
x > 3 & x <= 4
```

```
## [1] FALSE TRUE TRUE FALSE FALSE
```

```
which(x > 3 & x <= 4)
```

```
## [1] 2 3
```

```
all(x > 3)
```

```
## [1] FALSE
```

```
all(x > 3)
```

```
## [1] FALSE
```

```
any(x > 3)
```

```
## [1] TRUE
```

```
(1.5 - 0.5) == 1
```

```
## [1] TRUE
```

```
(1.9 - 0.9) == 1
```

```
## [1] FALSE
```

```
all.equal(1.9 - 0.9, 1)
```

```
## [1] TRUE
```

```
identical(1.5 - 0.5, 1)
```

```
## [1] TRUE
```

```
identical(1.9 - 0.9, 1)
```

```
## [1] FALSE
```

```
7 + TRUE
```

```
## [1] 8
```

```
7+FALSE
```

```
## [1] 7
```

Coercion

```
x
```

```
## [1] 1.80 3.14 4.00 88.17 13.00
```

```
is.numeric(x)
```

```
## [1] TRUE
```

```
is.character(x)
```

```
## [1] FALSE
```

```
as.character(x)
```

```
## [1] "1.8"      "3.14"     "4"        "88.169"  "13"
```

```
c(1, "a")
```

```
## [1] "1" "a"
```

Random number generation

```
set.seed(123)
```

```
rnorm(2)
```

```
## [1] -0.5605 -0.2302
```

```
rnorm(2)
```

```
## [1] 1.55871 0.07051
```

```
set.seed(123)
```

```
rnorm(2)
```

```
## [1] -0.5605 -0.2302
```


- Setting the seed to a specific value, simulations can be made exactly reproducible.

```
sample(1:5)
```

```
## [1] 5 1 2 3 4
```

```
sample(c("male", "female"), size=5, replace=TRUE,  
prob = c(0.2, 0.8))
```

```
## [1] "female" "male" "female" "female" "female"
```

Flow control

```
x <- c(1.8, 3.14, 4, 88.169, 13)
if (rnorm(1) > 0) sum(x) else mean(x)
## [1] 22.02

ifelse(x > 4, sqrt(x), x^2)
## [1] 3.240 9.860 16.000 9.390 3.606
```

```
x
```

```
## [1] 1.80 3.14 4.00 88.17 13.00
```

```
for(i in 2:5) {x[i]<-x[i]-x[i-1]}
```

```
x
```

```
## [1] 1.80 1.34 2.66 85.51 -72.51
```

```
x
```

```
## [1] 1.80 1.34 2.66 85.51 -72.51
```

```
while(sum(x)<100) {x<-2*x}
```

```
x
```

```
## [1] 14.40 10.72 21.28 684.07 -580.07
```

Writing functions

```
cmeans <- function(X)
{
  rval <- rep(0, ncol(X))
  for(j in 1:ncol(X))
  {
    mysum<-0
    for(i in 1:nrow(X)) mysum<-mysum+X[i,j]
    rval[j] <- mysum/ncol(X)
  }
  return(rval)
}
```

```
X <- matrix(1:20, ncol = 2)
X
```

```
##           [,1] [,2]
## [1,]         1  11
## [2,]         2  12
## [3,]         3  13
## [4,]         4  14
## [5,]         5  15
## [6,]         6  16
## [7,]         7  17
## [8,]         8  18
## [9,]         9  19
## [10,]        10  20
```

```
cmeans(X)
```

```
## [1]  5.5 15.5
```



```
colMeans(X)
```

```
## [1] 5.5 15.5
```

```
X <- matrix(rnorm(2*10^6), ncol = 2)
system.time(colMeans(X))
```

```
##      user  system elapsed
##    0.000    0.000    0.002
```

```
system.time(cmeans(X))
```

```
##      user  system elapsed
##    0.136    0.000    0.134
```

Vectorized calculations

```
cmeans2 <- function(X) {  
  rval <- rep(0, ncol(X))  
  for(j in 1:ncol(X)) rval[j] <- mean(X[,j])  
  return(rval)  
}
```

```
system.time(cmeans2(X))
```

```
##      user  system elapsed  
## 0.016   0.004   0.017
```

```
apply(X, 2, mean)
```

```
## [1] -0.0005206 -0.0015578
```

```
system.time(apply(X, 2, mean))
```

```
##      user  system elapsed  
##    0.020    0.004    0.027
```


- ① Element-wise computations should be avoided if vectorized computations are available.
- ② Optimized solutions (if available) typically perform better than the generic `for` or `apply()` solution.
- ③ Loops can be written more compactly using the `apply()` function.

Formulas

```
f<-y~x  
class(f)  
## [1] "formula"
```

- y is explained by x .


```
x<-seq(from = 0, to = 10, by = 0.5)
y<-2+3*x+rnorm(21)
plot(f)
```



```
lm(f)
```

```
##
```

```
## Call:
```

```
## lm(formula = f)
```

```
##
```

```
## Coefficients:
```

```
## (Intercept)          x
```

```
##          2.00          3.01
```

Data Management in R

Creation from scratch

```
mydata<-data.frame(one=1:10,two=11:20,three=21:30)
```

```
mydata
```

##	one	two	three
## 1	1	11	21
## 2	2	12	22
## 3	3	13	23
## 4	4	14	24
## 5	5	15	25
## 6	6	16	26
## 7	7	17	27
## 8	8	18	28
## 9	9	19	29
## 10	10	20	30

```
mydata<-as.data.frame(matrix(1:30,ncol=3))  
names(mydata)<-c("one", "two", "three")  
mydata
```

```
##      one two three  
## 1      1  11    21  
## 2      2  12    22  
## 3      3  13    23  
## 4      4  14    24  
## 5      5  15    25  
## 6      6  16    26  
## 7      7  17    27  
## 8      8  18    28  
## 9      9  19    29  
## 10     10  20    30
```

Subset selection

```
mydata$two
```

```
## [1] 11 12 13 14 15 16 17 18 19 20
```

```
mydata[, "two"]
```

```
## [1] 11 12 13 14 15 16 17 18 19 20
```

```
mydata[, 2]
```

```
## [1] 11 12 13 14 15 16 17 18 19 20
```

```
attach(mydata)
```

```
two
```

```
## [1] 11 12 13 14 15 16 17 18 19 20
```

```
detach(mydata)
```

```
with(mydata,two)
```

```
## [1] 11 12 13 14 15 16 17 18 19 20
```

```
mydata.sub<-subset(mydata,two <= 16,select=-two)
mydata.sub
```

```
##      one three
## 1      1     21
## 2      2     22
## 3      3     23
## 4      4     24
## 5      5     25
## 6      6     26
```


Import and export